

## Assignment 5 - GANs and VAEs for Data Augmentation

### Conditional Generative Adversarial Networks (GANs)

```
# imports
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow_docs.vis import embed
import matplotlib.pyplot as plt
import tensorflow as tf
import numpy as np
import imageio

# constants and hyperparameters
batch_size = 64
num_channels = 1
num_classes = 10
image_size = 28
latent_dim = 128

# Loading the MNIST dataset and preprocessing it
# We'll use all the available examples from both the training and test sets
(x_train, y_train), (x_test, y_test) =
keras.datasets.mnist.load_data()
all_digits = np.concatenate([x_train, x_test])
all_labels = np.concatenate([y_train, y_test])

# Scale the pixel values to [0, 1] range, add a channel dimension to
the images, and one-hot encode the labels
all_digits = all_digits.astype("float32") / 255.0
all_digits = np.reshape(all_digits, (-1, 28, 28, 1))
all_labels = keras.utils.to_categorical(all_labels, 10)

# Create tf.data.Dataset
dataset = tf.data.Dataset.from_tensor_slices((all_digits, all_labels))
dataset = dataset.shuffle(buffer_size=1024).batch(batch_size)

print(f"Shape of training images: {all_digits.shape}")
print(f"Shape of training labels: {all_labels.shape}")

Downloading data from https://storage.googleapis.com/tensorflow/tf-
keras-datasets/mnist.npz
11493376/11490434 [=====] - 1s 0us/step
11501568/11490434 [=====] - 1s 0us/step
Shape of training images: (70000, 28, 28, 1)
Shape of training labels: (70000, 10)
```

```

# Calculating the number of input channel for the generator and
discriminator
# In a regular (unconditional) GAN, we start by sampling noise (of
some fixed dimension) from a normal distribution
# In our case, (conditional GAN), we also need to account for the
class labels
# We will have to add the number of classes to the input channels of
the generator (noise input) as well as the discriminator (generated
image input)
generator_in_channels = latent_dim + num_classes
discriminator_in_channels = num_channels + num_classes
print("Generator in channels: ", generator_in_channels)
print("Discriminator in channels: ", discriminator_in_channels)

```

```

Generator in channels: 138
Discriminator in channels: 11

```

```

# Creating the discriminator and generator
# Create the discriminator.
discriminator = keras.Sequential(
    [
        keras.layers.InputLayer((28, 28, discriminator_in_channels)),
        layers.Conv2D(64, (3, 3), strides=(2, 2), padding="same"),
        layers.LeakyReLU(alpha=0.2),
        layers.Conv2D(128, (3, 3), strides=(2, 2), padding="same"),
        layers.LeakyReLU(alpha=0.2),
        layers.GlobalMaxPooling2D(),
        layers.Dense(1),
    ],
    name="discriminator",
)

```

```

# Create the generator.
generator = keras.Sequential(
    [
        keras.layers.InputLayer((generator_in_channels,)),
        # We want to generate 128 + num_classes coefficients to
reshape into a
# 7x7x(128 + num_classes) map.
        layers.Dense(7 * 7 * generator_in_channels),
        layers.LeakyReLU(alpha=0.2),
        layers.Reshape((7, 7, generator_in_channels)),
        layers.Conv2DTranspose(128, (4, 4), strides=(2, 2),
padding="same"),
        layers.LeakyReLU(alpha=0.2),
        layers.Conv2DTranspose(128, (4, 4), strides=(2, 2),
padding="same"),
        layers.LeakyReLU(alpha=0.2),
        layers.Conv2D(1, (7, 7), padding="same",
activation="sigmoid"),
    ]
)

```

```

],
name="generator",
)

# Creating a ConditionalGAN model
class ConditionalGAN(keras.Model):
    def __init__(self, discriminator, generator, latent_dim):
        super(ConditionalGAN, self).__init__()
        self.discriminator = discriminator
        self.generator = generator
        self.latent_dim = latent_dim
        self.gen_loss_tracker =
keras.metrics.Mean(name="generator_loss")
        self.disc_loss_tracker =
keras.metrics.Mean(name="discriminator_loss")

    @property
    def metrics(self):
        return [self.gen_loss_tracker, self.disc_loss_tracker]

    def compile(self, d_optimizer, g_optimizer, loss_fn):
        super(ConditionalGAN, self).compile()
        self.d_optimizer = d_optimizer
        self.g_optimizer = g_optimizer
        self.loss_fn = loss_fn

    def train_step(self, data):
        # Unpack the data.
        real_images, one_hot_labels = data

        # Add dummy dimensions to the labels so that they can be
concatenated with
# the images. This is for the discriminator.
        image_one_hot_labels = one_hot_labels[:, :, None, None]
        image_one_hot_labels = tf.repeat(
            image_one_hot_labels, repeats=[image_size * image_size]
        )
        image_one_hot_labels = tf.reshape(
            image_one_hot_labels, (-1, image_size, image_size,
num_classes)
        )

        # Sample random points in the latent space and concatenate the
labels.
        # This is for the generator.
        batch_size = tf.shape(real_images)[0]
        random_latent_vectors = tf.random.normal(shape=(batch_size,
self.latent_dim))
        random_vector_labels = tf.concat(
            [random_latent_vectors, one_hot_labels], axis=1

```

```

)

# Decode the noise (guided by labels) to fake images.
generated_images = self.generator(random_vector_labels)

# Combine them with real images. Note that we are
concatenating the labels
# with these images here.
fake_image_and_labels = tf.concat([generated_images,
image_one_hot_labels], -1)
real_image_and_labels = tf.concat([real_images,
image_one_hot_labels], -1)
combined_images = tf.concat(
    [fake_image_and_labels, real_image_and_labels], axis=0
)

# Assemble labels discriminating real from fake images.
labels = tf.concat(
    [tf.ones((batch_size, 1)), tf.zeros((batch_size, 1))],
axis=0
)

# Train the discriminator.
with tf.GradientTape() as tape:
    predictions = self.discriminator(combined_images)
    d_loss = self.loss_fn(labels, predictions)
    grads = tape.gradient(d_loss,
self.discriminator.trainable_weights)
    self.d_optimizer.apply_gradients(
        zip(grads, self.discriminator.trainable_weights)
    )

# Sample random points in the latent space.
random_latent_vectors = tf.random.normal(shape=(batch_size,
self.latent_dim))
random_vector_labels = tf.concat(
    [random_latent_vectors, one_hot_labels], axis=1
)

# Assemble labels that say "all real images".
misleading_labels = tf.zeros((batch_size, 1))

# Train the generator (note that we should *not* update the
weights
# of the discriminator)!
with tf.GradientTape() as tape:
    fake_images = self.generator(random_vector_labels)
    fake_image_and_labels = tf.concat([fake_images,
image_one_hot_labels], -1)

```

```

        predictions = self.discriminator(fake_image_and_labels)
        g_loss = self.loss_fn(misleading_labels, predictions)
        grads = tape.gradient(g_loss,
self.generator.trainable_weights)
        self.g_optimizer.apply_gradients(zip(grads,
self.generator.trainable_weights))

```

```

    # Monitor loss.
    self.gen_loss_tracker.update_state(g_loss)
    self.disc_loss_tracker.update_state(d_loss)
    return {
        "g_loss": self.gen_loss_tracker.result(),
        "d_loss": self.disc_loss_tracker.result(),
    }

```

*# Training the Conditional GAN*

```

cond_gan = ConditionalGAN(
    discriminator=discriminator, generator=generator,
    latent_dim=latent_dim
)
cond_gan.compile(
    d_optimizer=keras.optimizers.Adam(learning_rate=0.0003),
    g_optimizer=keras.optimizers.Adam(learning_rate=0.0003),
    loss_fn=keras.losses.BinaryCrossentropy(from_logits=True),
)

```

```
cond_gan.fit(dataset, epochs=20)
```

```

Epoch 1/20
1094/1094 [=====] - 950s 866ms/step - g_loss:
1.5757 - d_loss: 0.4204
Epoch 2/20
1094/1094 [=====] - 951s 869ms/step - g_loss:
1.2776 - d_loss: 0.4886
Epoch 3/20
1094/1094 [=====] - 928s 848ms/step - g_loss:
1.5587 - d_loss: 0.3958
Epoch 4/20
1094/1094 [=====] - 931s 851ms/step - g_loss:
2.3260 - d_loss: 0.1978
Epoch 5/20
1094/1094 [=====] - 949s 867ms/step - g_loss:
1.1429 - d_loss: 0.5997
Epoch 6/20
1094/1094 [=====] - 962s 879ms/step - g_loss:
0.9651 - d_loss: 0.6255
Epoch 7/20
1094/1094 [=====] - 889s 813ms/step - g_loss:
0.8912 - d_loss: 0.6428
Epoch 8/20

```

```

1094/1094 [=====] - 831s 759ms/step - g_loss:
0.8685 - d_loss: 0.6406
Epoch 9/20
1094/1094 [=====] - 834s 762ms/step - g_loss:
0.8420 - d_loss: 0.6571
Epoch 10/20
1094/1094 [=====] - 834s 762ms/step - g_loss:
0.7977 - d_loss: 0.6663
Epoch 11/20
1094/1094 [=====] - 834s 762ms/step - g_loss:
0.7840 - d_loss: 0.6755
Epoch 12/20
1094/1094 [=====] - 836s 764ms/step - g_loss:
0.7833 - d_loss: 0.6776
Epoch 13/20
1094/1094 [=====] - 834s 762ms/step - g_loss:
0.7547 - d_loss: 0.6828
Epoch 14/20
1094/1094 [=====] - 835s 763ms/step - g_loss:
0.7476 - d_loss: 0.6887
Epoch 15/20
1094/1094 [=====] - 834s 763ms/step - g_loss:
0.7496 - d_loss: 0.6816
Epoch 16/20
1094/1094 [=====] - 836s 765ms/step - g_loss:
0.7624 - d_loss: 0.6730
Epoch 17/20
1094/1094 [=====] - 834s 762ms/step - g_loss:
0.7621 - d_loss: 0.6702
Epoch 18/20
1094/1094 [=====] - 835s 763ms/step - g_loss:
0.7665 - d_loss: 0.6753
Epoch 19/20
1094/1094 [=====] - 834s 762ms/step - g_loss:
0.7768 - d_loss: 0.6621
Epoch 20/20
1094/1094 [=====] - 838s 766ms/step - g_loss:
0.7813 - d_loss: 0.6583

```

```
<keras.callbacks.History at 0x28124a223a0>
```

```

# Interpolating between classes with the trained generator
# We first extract the trained generator from our Conditiona GAN.
trained_gen = cond_gan.generator

```

```

# Choose the number of intermediate images that would be generated in
# between the interpolation + 2 (start and last images).
num_interpolation = 9 # @param {type:"integer"}

```

```
# Sample noise for the interpolation.
```

```

interpolation_noise = tf.random.normal(shape=(1, latent_dim))
interpolation_noise = tf.repeat(interpolation_noise,
                                repeats=num_interpolation)
interpolation_noise = tf.reshape(interpolation_noise,
                                (num_interpolation, latent_dim))

def interpolate_class(first_number, second_number):
    # Convert the start and end labels to one-hot encoded vectors.
    first_label = keras.utils.to_categorical([first_number],
num_classes)
    second_label = keras.utils.to_categorical([second_number],
num_classes)
    first_label = tf.cast(first_label, tf.float32)
    second_label = tf.cast(second_label, tf.float32)

    # Calculate the interpolation vector between the two labels.
    percent_second_label = tf.linspace(0, 1, num_interpolation)[:,
None]
    percent_second_label = tf.cast(percent_second_label, tf.float32)
    interpolation_labels = (
        first_label * (1 - percent_second_label) + second_label *
percent_second_label
    )

    # Combine the noise and the labels and run inference with the
generator.
    noise_and_labels = tf.concat([interpolation_noise,
interpolation_labels], 1)
    fake = trained_gen.predict(noise_and_labels)
    return fake

start_class = 1 # @param {type:"slider", min:0, max:9, step:1}
end_class = 5 # @param {type:"slider", min:0, max:9, step:1}

fake_images = interpolate_class(start_class, end_class)

# first sample noise from a normal distribution and then repeat that
for num_interpolation times and reshape the result accordingly
# then distribute it uniformly for num_interpolation with the label
identities being present in some proportion
fake_images *= 255.0
converted_images = fake_images.astype(np.uint8)
converted_images = tf.image.resize(converted_images, (96,
96)).numpy().astype(np.uint8)
imageio.mimsave("animation.gif", converted_images, fps=1)
embed.embed_file("animation.gif")

<IPython.core.display.HTML object>

```